

South Park Methodology

Kaylin Walker • February 2, 2016

Abstract

Quickly after Matt Stone and Trey Parker launched South Park, an adult animated series, in 1997, it became known for its use of crude language and dark satire of current events. Stone and Parker, the primary writers, producers and voice actors, have given the series a unique voice and personality. Each episode opens with the disclaimer:

“All characters and events in this show—even those based on real people—are entirely fictional. All celebrity voices are impersonated. . . .poorly. The following program contains coarse language and due to its content it should not be viewed by anyone.”

The series follows four main characters (Stan, Kyle, Cartman and Kenny), though it has an extensive ensemble cast of recurring characters. This analysis serves to analyze their speech to determine which words and phrases are distinct for each character. Since the series uses a lot of running gags, common phrases should be easy to find.

The programming language R and packages XML, RCurl, tm, RWeka and stringr were used to scrape South Park episode transcripts from the internet, attribute them to a certain character, break them into ngrams and identify characteristic ngrams for each character.

Method & Summary Statistics

Transcripts of the first 132 episodes of South Park (season 1 through season 9, episode 7) were scraped from unstructured text on the [Internet Movie Script Database](#) using the XML package. The series has run for 19 seasons and 267 episodes, but to my knowledge the remaining transcripts are not available online.

I was able to assign a speaker to each line by splitting the html at tags that contained only uppercase text and were shorter than 40 characters. Each line was followed by a blank line, so I used their index to create a starting and stopping point for text to be attributed to a speaker. From there, I used the tm package to pre-process the text (to lowercase, remove punctuation, numbers and whitespace; remove stopwords for unigrams and bigrams, but left them in for tri-, 4- and 5-grams) and form a corpus, which contained more than 18,000 unique words spoken more than 211,000 times. Reducing the sparsity brought that down to about 2,300 words spoken 172,000 times.

ngram.size	total	unique
1	172269	2349
2	140679	3026
3	19313	662
4	4925	360
5	1334	87

23 characters with the most words were retained, and the remaining 1781 speakers combined into one “all others” category so as not to lose the text.

speaker	words	speaker	words	speaker	words
cartman	28206	jimmy	2101	ms..cartman	1084
stan	19976	sharon	2041	ms..choksondik	980
kyle	17192	counselor.mackey	1955	kenny	931
butters	5731	announcer	1438	terrance	884
mr..garrison	5599	mayor	1163	reporter	866
chef	4836	chris	1137	mephesto	799
randy	4368	jesus	1120	narrator	789
jimbo	2767	wendy	1094	all.others	65212

Table 2: Number of Words by Character

Log Likelihood

Each corpus was analyzed to determine the most characteristic words for each speaker. Frequent and characteristic words are not the same thing - otherwise words like “I”, “school”, and “you” would rise to the top instead of unique words and phrases like “professor chaos”, “hippies” and “you killed kenny.”

Log likelihood was used to measure the unique-ness of the n-grams by character. Log likelihood compares the occurrence of a word in a particular corpus (the body of a character’s speech) to its occurrence in another corpus (all of the remaining South Park text) to determine if it shows up more or less likely than expected. The returned value represents the likelihood that the corpora are from the same, larger corpus, like a t-test. The higher the score, the more unlikely.

The **chi-square test** (χ^2), or goodness-of-fit test, can be used to compare the occurrence of a word across corpora.

$$\chi^2 = \sum \frac{(O_i - E_i)^2}{E_i}$$

where O = observed frequency and E = expected frequency.

However, flaws have been identified: invalidity at low frequencies (Dunning, 1993) and over-emphasis of common words (Kilgariff, 1996). Dunning was able to show that the **log-likelihood statistic** was accurate even at low frequencies:

$$2 \sum O_i * \ln\left(\frac{O_i}{E_i}\right)$$

Which can be computed from the contingency table below as $2 * ((a * \log(\frac{a}{E1}) + (b * \log(\frac{b}{E2})))$, where $E1 = (a + c) * \frac{(a+b)}{(c+d)}$, and $E2 = (b + d) * \frac{(a+b)}{(c+d)}$.

Group	Corpus.One	Corpus.Two	Total
Word	a	b	a+b
Not Word	c	d	c+d
Total	a+c	b+d	N=a+b+c+d

Table 3: Basic Framework

Group	Cartmans.Text	Remaining.Text	Total
‘hippies’	36	5	41
Not ‘hippies’	28170	144058	172228
Total	28206	144063	172269

Table 4: An Example with Log Likelihood 101.7

Computed:

$$E1 = 28206 * (41/172269) = 6.71 \quad E2 = 144063 * (41/172269) = 34.28$$

$$LL = 2 * [36 * \log(36/6.71) + 5 * \log(5/34.28)] = 101.7$$

Based on the overall ratio of the word “hippies” in the text, $41/172269 = 0.00023$, we would expect to see hippies in Cartman’s speech 6.71 times and in the remaining text 34.28 times. The log likelihood value of 101.7 is significant far beyond even the 0.01% level, so we can reject the null hypothesis that Cartman and the remaining text are one and the same.

For this analysis, a significance level of 0.001 was chosen to balance the number of significant ngrams with significance. 5.64% of ngrams were found to be significantly characteristic of a certain character.

Level	Critical.Value	P.Value	Percent.Sig
5%	3.84	0.05	19.93
1%	6.63	0.01	10.62
0.1%	10.83	.001	5.64
0.01%	15.13	.0001	3.47

Table 5: Log Likelihood Significance Levels

References

- Dunning, T. (1993) *Accurate Methods for the Statistics of Surprise and Coincidence*. Computational Linguistics, 19, 1, March 1993, pp. 61-74.
- Kilgariff, A. (1996) *Why chi-square doesn’t work, and an improved LOB-Brown comparison*. ALLC-ACH Conference, June 1996, Bergen, Norway.